

```
In [1]: ▶ import pandas as pd
import sqlite3
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
```

```
In [2]: ▶ df = pd.read_csv('bom.movie_gross.csv.')
df.head()
# df.shape
```

Out[2]:

	title	studio	domestic_gross	foreign_gross	year
0	Toy Story 3	BV	415000000.0	652000000	2010
1	Alice in Wonderland (2010)	BV	334200000.0	691300000	2010
2	Harry Potter and the Deathly Hallows Part 1	WB	296000000.0	664300000	2010
3	Inception	WB	292600000.0	535700000	2010
4	Shrek Forever After	P/DW	238700000.0	513900000	2010

```
In [3]: ▶ titles = df.title.unique()
len(titles)
### There are no repeats
```

Out[3]: 3386

```
In [4]: ▶ df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3387 entries, 0 to 3386
Data columns (total 5 columns):
#   Column                Non-Null Count  Dtype
---  -
0   title                  3387 non-null   object
1   studio                  3382 non-null   object
2   domestic_gross          3359 non-null   float64
3   foreign_gross            2037 non-null   object
4   year                    3387 non-null   int64
dtypes: float64(1), int64(1), object(3)
memory usage: 132.4+ KB
```

```
In [5]: df['foriegn_int'] = pd.to_numeric(df.foreign_gross.str.replace(',',''))

df['Total_Gross_Millions'] = (df.domestic_gross + df.foriegn_int) / 1000000
df.head()
```

Out[5]:

	title	studio	domestic_gross	foreign_gross	year	foriegn_int	Total_Gross_Milli
0	Toy Story 3	BV	415000000.0	652000000	2010	652000000.0	106
1	Alice in Wonderland (2010)	BV	334200000.0	691300000	2010	691300000.0	102
2	Harry Potter and the Deathly Hallows Part 1	WB	296000000.0	664300000	2010	664300000.0	96
3	Inception	WB	292600000.0	535700000	2010	535700000.0	82
4	Shrek Forever After	P/DW	238700000.0	513900000	2010	513900000.0	75

```
In [6]: df_moviebudgets = pd.read_csv('tn.movie_budgets.csv', index_col=0)
df_moviebudgets.head()
```

Out[6]:

	release_date	movie	production_budget	domestic_gross	worldwide_gross
id					
1	Dec 18, 2009	Avatar	\$425,000,000	\$760,507,625	\$2,776,345,279
2	May 20, 2011	Pirates of the Caribbean: On Stranger Tides	\$410,600,000	\$241,063,875	\$1,045,663,875
3	Jun 7, 2019	Dark Phoenix	\$350,000,000	\$42,762,350	\$149,762,350
4	May 1, 2015	Avengers: Age of Ultron	\$330,600,000	\$459,005,868	\$1,403,013,963
5	Dec 15, 2017	Star Wars Ep. VIII: The Last Jedi	\$317,000,000	\$620,181,382	\$1,316,721,747

```
In [7]: df_moviebudgets.shape
```

Out[7]: (5782, 5)

```
In [8]: movie = df_moviebudgets.movie.unique()
len(movie)
### There are a few duplicates
```

Out[8]: 5698

```
In [9]: # df_reviews = pd.read_csv('rt.reviews.tsv', sep = '\t')
df_movieinfo = pd.read_table('rt.movie_info.tsv', index_col=0)
df_movieinfo.head()
```

Out[9]:

	synopsis	rating	genre	director	writer	theater_date
id						
1	This gritty, fast-paced, and innovative police...	R	Action and Adventure Classics Drama	William Friedkin	Ernest Tidyman	Oct 9, 1971
3	New York City, not-too-distant-future: Eric Pa...	R	Drama Science Fiction and Fantasy	David Cronenberg	David Cronenberg Don DeLillo	Aug 17, 2012
5	Illeana Douglas delivers a superb performance ...	R	Drama Musical and Performing Arts	Allison Anders	Allison Anders	Sep 13, 1996
6	Michael Douglas runs afoul of a treacherous su...	R	Drama Mystery and Suspense	Barry Levinson	Paul Attanasio Michael Crichton	Dec 9, 1994
7	NaN	NR	Drama Romance	Rodney Bennett	Giles Cooper	NaN

```
In [10]: ### Connect to the IMDB Database
conn = sqlite3.connect('im.db')
```

```
In [11]: q = """
SELECT *
FROM movie_basics
JOIN movie_ratings USING(movie_id)
;
"""

x = pd.read_sql(q, conn)
pd.read_sql(q, conn)
```

Out[11]:

	movie_id	primary_title	original_title	start_year	runtime_minutes	ge
0	tt0063540	Sunghursh	Sunghursh	2013	175.0	Action,Crime,Di
1	tt0066787	One Day Before the Rainy Season	Ashad Ka Ek Din	2019	114.0	Biography,Di
2	tt0069049	The Other Side of the Wind	The Other Side of the Wind	2018	122.0	Di
3	tt0069204	Sabse Bada Sukh	Sabse Bada Sukh	2018	NaN	Comedy,Di
4	tt0100275	The Wandering Soap Opera	La Telenovela Errante	2017	80.0	Comedy,Drama,Far
...	
73851	tt9913084	Diabolik sono io	Diabolik sono io	2019	75.0	Docume
73852	tt9914286	Sokagin Çocuklari	Sokagin Çocuklari	2019	98.0	Drama,Fa
73853	tt9914642	Albatross	Albatross	2017	NaN	Docume
73854	tt9914942	La vida sense la Sara Amat	La vida sense la Sara Amat	2019	NaN	t
73855	tt9916160	Drømmeland	Drømmeland	2019	72.0	Docume

73856 rows × 8 columns

```
In [12]: df_budgets = pd.read_csv('tn.movie_budgets.csv', index_col=0)
df_budgets.head()
```

Out[12]:

	release_date	movie	production_budget	domestic_gross	worldwide_gross
id					
1	Dec 18, 2009	Avatar	\$425,000,000	\$760,507,625	\$2,776,345,279
2	May 20, 2011	Pirates of the Caribbean: On Stranger Tides	\$410,600,000	\$241,063,875	\$1,045,663,875
3	Jun 7, 2019	Dark Phoenix	\$350,000,000	\$42,762,350	\$149,762,350
4	May 1, 2015	Avengers: Age of Ultron	\$330,600,000	\$459,005,868	\$1,403,013,963
5	Dec 15, 2017	Star Wars Ep. VIII: The Last Jedi	\$317,000,000	\$620,181,382	\$1,316,721,747

```
In [13]: df_budgets['budget_millions'] = df_budgets.production_budget.str.replace(
df_budgets['budget_millions'] = df_budgets.budget_millions.str.replace('$')
df_budgets['budget_millions'] = pd.to_numeric(df_budgets.budget_millions)
df_budgets.head()
```

C:\Users\deapillai\AppData\Local\Temp\ipykernel_38760\2653470430.py:2: FutureWarning: The default value of regex will change from True to False in a future version. In addition, single character regular expressions will *not* be treated as literal strings when regex=True.

```
df_budgets['budget_millions'] = df_budgets.budget_millions.str.replace('$', '')
```

Out[13]:

	release_date	movie	production_budget	domestic_gross	worldwide_gross	budget_
id						
1	Dec 18, 2009	Avatar	\$425,000,000	\$760,507,625	\$2,776,345,279	
2	May 20, 2011	Pirates of the Caribbean: On Stranger Tides	\$410,600,000	\$241,063,875	\$1,045,663,875	
3	Jun 7, 2019	Dark Phoenix	\$350,000,000	\$42,762,350	\$149,762,350	
4	May 1, 2015	Avengers: Age of Ultron	\$330,600,000	\$459,005,868	\$1,403,013,963	
5	Dec 15, 2017	Star Wars Ep. VIII: The Last Jedi	\$317,000,000	\$620,181,382	\$1,316,721,747	

In [14]:

Create Dataframe

```
df_runtime_minutes = pd.merge(x,df_budgets,how='inner',left_on = 'primary_
df_runtime_minutes.head()
```

Out[14]:

	movie_id	primary_title	original_title	start_year	runtime_minutes	gen
0	tt0249516	Foodfight!	Foodfight!	2012	91.0	Action,Animation,Com
1	tt0326592	The Overnight	The Overnight	2010	88.0	No
2	tt3844362	The Overnight	The Overnight	2015	79.0	Comedy,Mysl
3	tt0337692	On the Road	On the Road	2012	124.0	Adventure,Drama,Roma
4	tt4339118	On the Road	On the Road	2014	89.0	Dra

In [15]:

```
df_runtime = df_runtime_minutes.drop_duplicates(subset=['movie'],
                                                inplace=False)
df_runtime.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 2126 entries, 0 to 2874
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  -
0   movie_id              2126 non-null   object
1   primary_title         2126 non-null   object
2   original_title        2126 non-null   object
3   start_year            2126 non-null   int64
4   runtime_minutes       2072 non-null   float64
5   genres                2124 non-null   object
6   averagerating         2126 non-null   float64
7   numvotes              2126 non-null   int64
8   release_date          2126 non-null   object
9   movie                 2126 non-null   object
10  production_budget     2126 non-null   object
11  domestic_gross        2126 non-null   object
12  worldwide_gross       2126 non-null   object
13  budget_millions       2126 non-null   float64
dtypes: float64(3), int64(2), object(9)
memory usage: 249.1+ KB
```

In [16]:

```
import seaborn as sns
import scipy.stats as stats
```

In [17]: `df_runtime.describe()`

Out[17]:

	start_year	runtime_minutes	averagerating	numvotes	budget_millions
count	2126.000000	2072.000000	2126.000000	2.126000e+03	2126.000000
mean	2013.561148	103.881757	6.231232	8.212615e+04	37.163874
std	2.512421	18.760685	1.142910	1.465604e+05	51.129898
min	2010.000000	5.000000	1.600000	5.000000e+00	0.001400
25%	2011.000000	91.000000	5.600000	7.930000e+02	5.000000
50%	2013.000000	102.000000	6.300000	2.344900e+04	18.250000
75%	2015.000000	114.000000	7.000000	9.486275e+04	45.000000
max	2019.000000	180.000000	9.200000	1.841066e+06	425.000000

In [18]: `pd.options.mode.chained_assignment = None`

In [19]: `#Domestic`
`df_runtime['domestic_millions'] = df_runtime.domestic_gross.str.replace('$', '')`
`df_runtime['domestic_millions'] = df_runtime.domestic_millions.str.replace('$', '')`
`df_runtime['domestic_millions'] = pd.to_numeric(df_runtime.domestic_millions, errors='coerce')`
`# # Worldwide`
`df_runtime['worldwide_millions'] = df_runtime.worldwide_gross.str.replace('$', '')`
`df_runtime['worldwide_millions'] = df_runtime.worldwide_millions.str.replace('$', '')`
`df_runtime['worldwide_millions'] = pd.to_numeric(df_runtime.worldwide_millions, errors='coerce')`

C:\Users\deapillai\AppData\Local\Temp\ipykernel_38760\2282399076.py:3: FutureWarning: The default value of regex will change from True to False in a future version. In addition, single character regular expressions will *not* be treated as literal strings when regex=True.

`df_runtime['domestic_millions'] = df_runtime.domestic_millions.str.replace('$', '')`

C:\Users\deapillai\AppData\Local\Temp\ipykernel_38760\2282399076.py:7: FutureWarning: The default value of regex will change from True to False in a future version. In addition, single character regular expressions will *not* be treated as literal strings when regex=True.

`df_runtime['worldwide_millions'] = df_runtime.worldwide_millions.str.replace('$', '')`

In [20]: `df_runtime['domestic_millions'] = pd.to_numeric(df_runtime.domestic_millions, errors='coerce')`
`df_runtime['worldwide_millions'] = pd.to_numeric(df_runtime.worldwide_millions, errors='coerce')`

```
In [21]: df_runtime['production_budget'] = df_runtime.production_budget.str.replace(
df_runtime['production_budget'] = df_runtime.production_budget.str.replace(
df_runtime['production_budget'] = pd.to_numeric(df_runtime.production_budg
```

C:\Users\deapillai\AppData\Local\Temp\ipykernel_38760\1529873331.py:2: FutureWarning: The default value of regex will change from True to False in a future version. In addition, single character regular expressions will *not* be treated as literal strings when regex=True.

```
df_runtime['production_budget'] = df_runtime.production_budget.str.replace('$', '')
```

```
In [22]: df_runtime['production_budget'] = pd.to_numeric(df_runtime.production_budg
```

```
In [23]: df_runtime['ROI'] = (df_runtime.worldwide_millions - df_runtime.production
```

```
In [24]: df_runtime.ROI.describe()
```

```
Out[24]: count    2126.000000
mean         2.627956
std          13.627146
min          -1.000000
25%          -0.635096
50%           0.685090
75%           2.644018
max          415.564740
Name: ROI, dtype: float64
```

```
In [25]: df_runtime.head()
```

Out[25]:

	movie_id	primary_title	original_title	start_year	runtime_minutes	gen
0	tt0249516	Foodfight!	Foodfight!	2012	91.0	Action,Animation,Com
1	tt0326592	The Overnight	The Overnight	2010	88.0	No
3	tt0337692	On the Road	On the Road	2012	124.0	Adventure,Drama,Roma
6	tt0359950	The Secret Life of Walter Mitty	The Secret Life of Walter Mitty	2013	114.0	Adventure,Comedy,Dra
7	tt0365907	A Walk Among the Tombstones	A Walk Among the Tombstones	2014	114.0	Action,Crime,Dra

```
In [26]: df_runtime['runtime_minutes'].mean()
```

Out[26]: 103.88175675675676


In [27]: `import seaborn as sns`

In [28]: `df_runtime.rename(columns={"production_budget":"budget_millions"})`

Out[28]:

	movie_id	primary_title	original_title	start_year	runtime_minutes	
0	tt0249516	Foodfight!	Foodfight!	2012	91.0	Action,Animation,C
1	tt0326592	The Overnight	The Overnight	2010	88.0	
3	tt0337692	On the Road	On the Road	2012	124.0	Adventure,Drama,Rc
6	tt0359950	The Secret Life of Walter Mitty	The Secret Life of Walter Mitty	2013	114.0	Adventure,Comedy
7	tt0365907	A Walk Among the Tombstones	A Walk Among the Tombstones	2014	114.0	Action,Crime
...	
2870	tt8680254	Richard III	Richard III	2016	NaN	
2871	tt8824064	Heroes	Heroes	2019	88.0	Docun
2872	tt8976772	Push	Push	2019	92.0	Docun
2873	tt9024106	Unplanned	Unplanned	2019	106.0	Biography
2874	tt9248762	The Terrorist	The Terrorist	2018	NaN	


2126 rows × 7 columns



In [29]: `df_runtime.head()`


Out[29]:

	movie_id	primary_title	original_title	start_year	runtime_minutes	gen
0	tt0249516	Foodfight!	Foodfight!	2012	91.0	Action,Animation,Com
1	tt0326592	The Overnight	The Overnight	2010	88.0	No
3	tt0337692	On the Road	On the Road	2012	124.0	Adventure,Drama,Roma
6	tt0359950	The Secret Life of Walter Mitty	The Secret Life of Walter Mitty	2013	114.0	Adventure,Comedy,Dra
7	tt0365907	A Walk Among the Tombstones	A Walk Among the Tombstones	2014	114.0	Action,Crime,Dra




```
In [30]: sns.histplot(x=df_runtime.runtime_minutes,  
                      y=df_runtime.ROI)
```

```
Out[30]: <AxesSubplot:xlabel='runtime_minutes', ylabel='ROI'>
```




```
In [31]: sns.lineplot(x=df_runtime.runtime_minutes,  
                      y=df_runtime.ROI)
```

```
Out[31]: <AxesSubplot:xlabel='runtime_minutes', ylabel='ROI'>
```



```
In [32]: sns.scatterplot(x=df_runtime.runtime_minutes,  
                        y=df_runtime.ROI)
```

```
Out[32]: <AxesSubplot:xlabel='runtime_minutes', ylabel='ROI'>
```




```
In [33]: df = df_runtime.drop(df_runtime[df_runtime['averagerating'] < 7.1].index,
df.info())
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 515 entries, 1 to 2872
Data columns (total 17 columns):
#   Column                Non-Null Count  Dtype
---  -
0   movie_id              515 non-null   object
1   primary_title         515 non-null   object
2   original_title        515 non-null   object
3   start_year            515 non-null   int64
4   runtime_minutes       496 non-null   float64
5   genres                514 non-null   object
6   averagerating         515 non-null   float64
7   numvotes              515 non-null   int64
8   release_date          515 non-null   object
9   movie                 515 non-null   object
10  production_budget     515 non-null   float64
11  domestic_gross        515 non-null   object
12  worldwide_gross       515 non-null   object
13  budget_millions       515 non-null   float64
14  domestic_millions     515 non-null   float64
15  worldwide_millions    515 non-null   float64
16  ROI                   515 non-null   float64
dtypes: float64(7), int64(2), object(8)
memory usage: 72.4+ KB
```


```
In [34]: sns.scatterplot(x=df.runtime_minutes,  
                        y=df.averagerating)
```

```
Out[34]: <AxesSubplot:xlabel='runtime_minutes', ylabel='averagerating'>
```




```
In [35]: sns.lineplot(x=df.runtime_minutes,  
                     y=df.averagerating)
```

```
Out[35]: <AxesSubplot:xlabel='runtime_minutes', ylabel='averagerating'>
```




```
In [36]: sns.histplot(x=df.runtime_minutes,  
                      y=df.averagerating)
```

```
Out[36]: <AxesSubplot:xlabel='runtime_minutes', ylabel='averagerating'>
```



```
In [37]: sns.histplot(x=df.averagerating  
)
```

```
Out[37]: <AxesSubplot:xlabel='averagerating', ylabel='Count'>
```




```
In [38]: df.averagerating.describe()
```

```
Out[38]: count      515.000000  
mean         7.574951  
std          0.438752  
min          7.100000  
25%          7.200000  
50%          7.400000  
75%          7.800000  
max          9.200000  
Name: averagerating, dtype: float64
```




```
In [39]: sns.lineplot(x=df_runtime.runtime_minutes,  
                      y=df_runtime.averagerating)
```

```
Out[39]: <AxesSubplot:xlabel='runtime_minutes', ylabel='averagerating'>
```




```
In [40]: sns.histplot(x=df.runtime_minutes,  
                      y=df.averagerating)
```

```
Out[40]: <AxesSubplot:xlabel='runtime_minutes', ylabel='averagerating'>
```



```
In [41]: sns.histplot(x=df_runtime.runtime_minutes, kde=True)
```

```
Out[41]: <AxesSubplot:xlabel='runtime_minutes', ylabel='Count'>
```




```
In [42]: df.describe()
```

```
Out[42]:
```

	start_year	runtime_minutes	averagerating	numvotes	production_budget	bud
count	515.000000	496.000000	515.000000	5.150000e+02	515.000000	
mean	2013.852427	110.516129	7.574951	1.729716e+05	45.053837	
std	2.537188	23.781831	0.438752	2.341522e+05	59.366843	
min	2010.000000	5.000000	7.100000	5.000000e+00	0.015000	
25%	2012.000000	96.000000	7.200000	9.765000e+02	8.000000	
50%	2014.000000	109.000000	7.400000	8.157500e+04	20.000000	
75%	2016.000000	126.250000	7.800000	2.630700e+05	51.350000	
max	2019.000000	180.000000	9.200000	1.841066e+06	330.600000	

```
In [43]: sns.histplot(x=df.runtime_minutes, kde=True,)
```

```
Out[43]: <AxesSubplot:xlabel='runtime_minutes', ylabel='Count'>
```



```
In [44]: df2 = df_runtime.drop(df_runtime[df_runtime['averagerating'] > 7.1].index)
df2.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 1680 entries, 0 to 2874
Data columns (total 17 columns):
#   Column                Non-Null Count  Dtype
---  -
0   movie_id              1680 non-null   object
1   primary_title         1680 non-null   object
2   original_title        1680 non-null   object
3   start_year            1680 non-null   int64
4   runtime_minutes       1645 non-null   float64
5   genres                1679 non-null   object
6   averagerating         1680 non-null   float64
7   numvotes              1680 non-null   int64
8   release_date          1680 non-null   object
9   movie                 1680 non-null   object
10  production_budget     1680 non-null   float64
11  domestic_gross        1680 non-null   object
12  worldwide_gross       1680 non-null   object
13  budget_millions       1680 non-null   float64
14  domestic_millions     1680 non-null   float64
15  worldwide_millions    1680 non-null   float64
16  ROI                   1680 non-null   float64
dtypes: float64(7), int64(2), object(8)
memory usage: 236.2+ KB
```

```
In [46]: sns.histplot(x=df2.runtime_minutes, bins = 10, kde=True,)
```

```
Out[46]: <AxesSubplot:xlabel='runtime_minutes', ylabel='Count'>
```

